

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [ ]: #Load the dataset
df = pd.read_csv("/content/Salary_dataset.csv")
```

```
In [ ]: data = pd.DataFrame(df)
```

```
In [ ]: df.head()
```

```
Out[ ]:      Unnamed: 0  YearsExperience  Salary
0              0              1.2  39344.0
1              1              1.4  46206.0
2              2              1.6  37732.0
3              3              2.1  43526.0
4              4              2.3  39892.0
```

```
In [ ]: df.tail()
```

```
Out[ ]:      Unnamed: 0  YearsExperience  Salary
25              25              9.1  105583.0
26              26              9.6  116970.0
27              27              9.7  112636.0
28              28             10.4  122392.0
29              29             10.6  121873.0
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      30 non-null    int64
1   YearsExperience 30 non-null    float64
2   Salary          30 non-null    float64
dtypes: float64(2), int64(1)
memory usage: 852.0 bytes
```

```
In [ ]: data.dropna(inplace=True)
```

```
In [ ]: data.drop_duplicates(inplace=True)
```

```
In [ ]: def remove_outliers(df, column):  
        from scipy import stats  
        z_scores = np.abs(stats.zscore(df[column]))  
        return df[(z_scores < 3)]  
  
data = remove_outliers(data, 'Salary')
```

```
In [ ]: data = data.astype({'YearsExperience': 'float', 'Salary': 'int'})
```

```
In [ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 30 entries, 0 to 29  
Data columns (total 3 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   Unnamed: 0      30 non-null    int64  
1   YearsExperience 30 non-null    float64  
2   Salary          30 non-null    int64  
dtypes: float64(1), int64(2)  
memory usage: 852.0 bytes
```

```
In [ ]: # Visualizing the data before training  
plt.scatter(data['YearsExperience'], data['Salary'], color='blue', label='Ac  
plt.xlabel("Years of Experience")  
plt.ylabel("Salary")  
plt.title("Salary vs Experience (Before Model Training)")  
plt.legend()  
plt.show()
```



```
In [ ]: #Splitting data into training and testing sets
X = data[['YearsExperience']]
y = data['Salary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

```
In [ ]: #Train the model
model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[ ]: ▼ LinearRegression ⓘ ?
LinearRegression()
```

```
In [ ]: #Predictions on test set
y_pred = model.predict(X_test)
```

```
In [ ]: #Visualizing the regression line after training
plt.scatter(X, y, color='blue', label='Actual Data')
plt.plot(X, model.predict(X), color='black', label='Regression Line')
plt.xlabel("Years of Experience")
plt.ylabel("Salary")
plt.title("Salary vs Experience (After Model Training)")
plt.legend()
plt.show()
```




```
In [ ]: #Model Evaluation
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```


```
In [ ]: print(f"Mean Squared Error: {mse}")
print(f"Mean Absolute Error: {mae}")
print(f"R2 Score: {r2}")
```

Mean Squared Error: 49830096.855908394  
Mean Absolute Error: 6286.453830757745  
R<sup>2</sup> Score: 0.9024461774180497


```
In [ ]: plt.scatter(y_test, y_pred, color='blue')
plt.xlabel("Actual Salary")
plt.ylabel("Predicted Salary")
plt.title("Actual vs. Predicted Salary")
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='black')
plt.show()
```



```
In [ ]: #Residual Plot (Error Analysis)
residuals = y_test - y_pred
plt.scatter(y_test, residuals, color='blue')
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel("Actual Salary")
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()
```



```
In [ ]: #Histogram of Residuals
sns.histplot(residuals, kde=True, bins=15, color="green")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.title("Distribution of Residuals")
plt.show()
```



```
In [ ]: #Validation with sample data
sample_experience = np.array([3, 4.5, 5.2]).reshape(-1, 1)
sample_experience = pd.DataFrame({'YearsExperience': [3, 4.5, 5.2]})
predicted_salaries = model.predict(sample_experience)
```

```
In [ ]: #Display results
for exp, sal in zip(sample_experience.values.flatten(), predicted_salaries):
    print(f"Years of Experience: {exp}, Predicted Salary: {sal:.2f}")
```

Years of Experience: 3.0, Predicted Salary: 52651.65  
Years of Experience: 4.5, Predicted Salary: 66787.37  
Years of Experience: 5.2, Predicted Salary: 73384.04